

# Comparing subtelomeric sequences in human genomes in terms of sequence similarity

525183, Jindřich Matuška

## Abstract

Subtelomeres are parts of DNA close to the telomeres. They are defined as segments between telomeric sequences and chromatin, in humans they are usually considered 500Kbp from each telomere. Subtelomeres are sequences with high similarity. They contain autonomously replicating sequence, tandem repeats and more. This allows for change in length of individual chromosomes, creation of space for de-novo genes etc.

In this project I explore similarity between individual subtelomeres. First I create a pipeline for easy and efficient extraction of subtelomeres. Then I look into possibilities of similarity analysis using tool ModDotPlot with the help of custom script for generating heatmaps from resulting `.bedpe` file. Finally I introduce a method for searching for frequent sequences using combination of approaches such as markov model using graph of kmers and great deluge.

Each part of this report contains a sample script, which may be used for running the tools and pipelines. Created tools are also documented both using `-h` option and inside scripts.

## Project repository: getting started

This project has an corresponding repository under following link:

[github.com/Ardnij123/pv269\\_project](https://github.com/Ardnij123/pv269_project)

To prevent the repository being too large, the analysed sequence has not been uploaded to it. You can either use your own sequences or download one. I have worked with sequence HG002v1.1 [1]. The sequence should be a telomere-to-telomere sequence in high quality. For use with `repeats.py`, it is preferred to not contain any N's. Similarly, most of intermediate results are not uploaded to the repository.

Most of the tools used can be downloaded using Conda. The list as an installation script can be found at `.conda_files/conda_create.sh` [2-5]. Other than that, I have used package ModDotPlot [6], which may be either directly installed or used for instance as Docker image, for instance this one.

## Subtelomere extraction

```
# Extract subtelomeres of size 500Kbp in file hg002v1.1.fa.gz  
# and write it into file subtelomeres.fa  
./work/get_subtelomeres/extract_subtelomeres.sh -n 500000 -o subtelomeres.fa
```

I have created a pipeline for the extraction of subtelomeres. The pipeline consists of filtering out telomeric sequences, cropping, reversing end sequences and optionally extracting and rebasing a bedfile to correspond to the sequence names in the resulting fasta file. The script is parametrized to allow for extraction of any-length telomere sequence or use existing file with sizes of chromosomes for better performance.

Full pipeline can be found at `work/get_subtelomeres/extract_subtelomeres.sh`. Tools such as SeqKit [2], BEDTools [3] or Seqtk [4] are leveraged for most of the work. Scripts in Python or Bash are used for the remainder of steps.

The pipeline is quite fast (at least compared to other steps in this project). By far the slowest part is gaining chromosome sizes. One feature it lacks is that it does not check if there is any telomere sequence. This may lead to addition of unwanted sequence from not-the-subtelomere into the output.

## Graph of similiarity using ModDotPlot, heatmap and groups

```
mkdir plot-full
```

```
# Merge sequences
```

```
./work/similarity-plot/merge_split.sh subtelomeres.fa -o plot-full/merged.fa
```

```
# Run ModDotPlot in low resolution (window 5Kbp)
```

```
moddotplot static -f merged.fa -o plot-full/w5000 -w 5000 --no-plot --no-hist
```

```
# Visualise the bedpe file
```

```
./work/similarity-plot/event_counter.py plot-full/w5000/merged.bedpe \
    -s plot-full/merged.fa.seqs
```

```
# Select promissing sequences and plot in high res (window 1Kbp)
```

```
mkdir group_1
```

```
merge_split.sh subtelomeres.fa \
    -p chr1_*_START -p chr4_*_START -p chr7_*_END \
    -r "s/^>(*)_END/>Z_1/" -r "s/^>(*)_START/>A_1/" \
    -o group1/merged.fa
```

```
moddotplot static -f group1/merged.fa -o group1/w1000 -w 1000
```

For exploration of similiarity I have first used ModDotPlot [6]. ModDotPlot takes several fasta files on the input and approximates the average nucleotide identity between pairwise combinations of genomic intervals (tldr. it creates a hit heatmap for the files). I have written some notes on running ModDotPlot into `work/similarity-plot/moddotplot.md`.

ModDotPlot seems to have options for plotting allignment of multiple sequences in a grid. I have however not found a way to make it grid more than 4 sequences at once. For this I have created tiny script that merges multiple sequences into

one, splitting them with a sequence of N's. This way, the graph generated by ModDotPlot is more readable.

The ModDotPlot seems to scale with the number of genomic intervals. It is best to first create a plot with low resolution, then choose some subset of sequences that seems to be promising (eg. subset with high number of similar intervals) and only create plot in higher resolution for this subset.




Figure 1: Output from ModDotPlot on all paternal chromosomes with window size 5000 bases. Only paternal copies are plotted and shown

At least in the first alignment of all subtelomeres, the resulting plot might be quite large and hard to comprehend (see Figure 1). To make the reading of it easier, I have created a Python script that plots the identities as a heatmap using Matplotlib. It also allows for turning on/off some sequences to better see the groups (see Figure 2). The script can be found at [work/similarity-plot/event\\_counter.py](#)




Figure 2: Output from ModDotPlot visualized using Event counter, same as in Figure 1. The order of subtelomeres is different. White means higher number of hits. Scale is logarithmic.

With output plotted in this way it is up to the human behind the computer to choose some groups of sequences for further examination. In this case, it seems to me that there are 2 groups of high similarity, that is:

- 1) starts of chromosomes 13, 14, 15 and 22 (see Figure 3)
- 2) starts of 16, X, Y and ends of 8, 12, 16, 19, 20, 22 (see Figure 4)

Other than that, start of the chromosome 21 has a high self-similarity (see Figure 5).




Figure 3: Output from ModDotPlot for group 1) with window size 1000 bases.

## Aligning sequences

ModDotPlot outputs besides the images a `.bedpe` file, which contains the hits in a format similar to the `bed` format, which makes extraction of the hit sequences from the full DNA file easy. I have tried to align the hits using ClustalO [7],




Figure 4: Output from ModDotPlot for group 2) with window size 1000 bases.

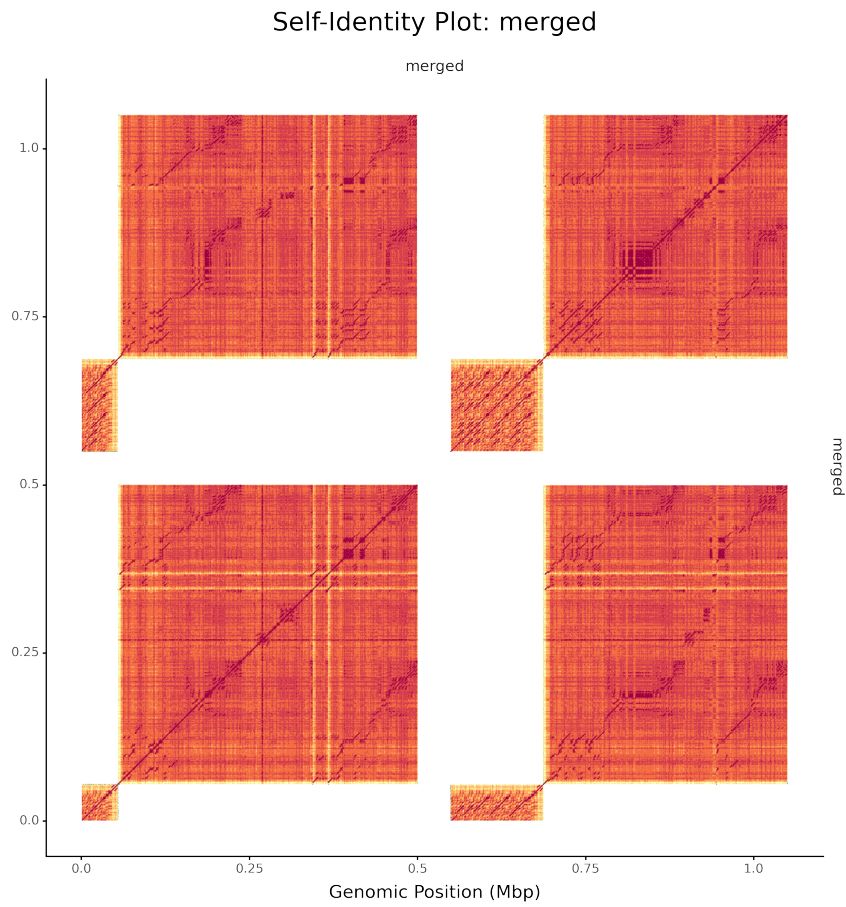


Figure 5: Output from ModDotPlot for start of chromosome 21 with window size 1000 bases.

however it ran very slow and in the end it failed. My assumption is that it was due to the very large number of short sequences. With better choice of aligner I would assume this could lead to quite easy extraction of the backbone of repeating sequences.

For script, see `work/repeat-search/alignment.sh`

## Sequences with high frequency

```
./work/repeat-search/repeats.py subtelomeres.fa \
    -k 20 -t 5 -i 10 -g 10 -b 3 -s 0 -e 12 -m 500 > repeats.bed
```

Finally, I have implemented my own algorithm for seeking subsequences with high frequency. This algorithm is based on markov chain and great deluge algorithm. First, it creates a graph with kmers as nodes and counts of transitions by reading single base as weights of edges. This graph is then scaled logarithmically and pruned for better performance. Weights then correspond to increments of score gained by passing through the edge (by reading the corresponding base). In the second run, the graph is used to find the path in it with the highest score. With some penalties gaps and insertions are allowed to model possible modifications to the DNA sequence.

Init:

Graph  $G$ ,  $V(G)$  - states of automaton,  $E(G)$  allowed transitions by reading sequence  
 weights of edges correspond to increment of value by moving along the edge  
 Sequence corresponding to the graph

Parameters:

MaxDrop, InsertionPenalty, GapPenalty, BasePenalty

```
1) Values[0] := {(v: 0) for v in V(G)}; Starts[0] := {(v: 0) for v in V(G)};
   MaxValue := 0; MaxPosition = 0; MinPosition = 0; Position = 0
```

```
2) Do {
```

```
    Position := Position + 1
```

```
    Base := following base on input
```

```
    Set Value[Position][v] to maximum of:
```

```
        ( Value[Position-1][v] - InsertionPenalty )
```

```
        ( Value[Position-1][v-1] - GapPenalty )
```

```
            if v-1 is node with edge coming to v
```

```
        ( Value[Position-1][v-1] + weight of the transition edge )
```

```
            if v-1 is node with edge coming to v corresponding to Base
```

```
Values[Position] := Values[Position] - BasePenalty
```

```
Set Starts[Position] according to the states
```



```

    For each State v s.t. Values[Position][v] <= 0; Do
        Values[Position][v] = 0
        Starts[v] = Position
    Done

    Score := max(Values)
    If MaxValue < Score; Then
        MaxValue := Score
        MaxPosition := Position
        MinPosition := Starts[argmax(Values)]
    Fi

    Optionally: Remove states v from Values where Values[v] < MaxValue - MaxDrop

} While Score >= MaxValue - MaxDrop

Return MinPosition, MaxPosition, MaxValue

```

For example with  $k=3$  and sequence AATTAATAATTA, there are 5 distinct kmers AAT, ATT, TTA, TAA, ATA. In the graph, the counts on the edges would be:

```

AAT - ATT - 2
ATT - TTA - 2
TTA - TAA - 1
TAA - AAT - 2
AAT - ATA - 1
ATA - TAA - 1

```

The values would then be rescaled using logarithmic scale, and the low ones would be removed (e.g. edges with count 1).

The search algorithm for the second run then reads bases one by one and in each step it computes best possible values for all nodes (kmers) in the graph, down to some minimal value (it is called `_flood` in the implementation). In each step and node, it also has 3 possible things to do. Either it reads a correct base, in which case it simply adds the weight of the edge to the previous value of node and uses it as new value of the next node. Or there might have been an insertion into the DNA, which means that it subtracts the insertion penalty of the node and uses it as a new value for the node. The third case could be that there was some deletion from the DNA, which corresponds to transition in the graph without reading base.

The implementation contains many optimisations such as starting only on exact match of given length, evaluation of nodes in order by descending score to not expand one node multiple times or using arrays for values over and over to mitigate reinitialization of them for each step. The results are more or less stable with respect to the parameters.

In the current state, it however does only print out the subsequences with

frequent kmers, not the corresponding probable sequence before mutations. It should not be hard to implement, but it is too close to the deadline. Also, it should be possible to drastically lower the memory consumption when creating the graph of kmers. (Currently it takes close to 16GB of memory for 72Mbp fasta file, but after the pruning and change of representation it goes below 2GB.)

Program prints out sequence of lines similar to **bed** format (except there is score in 4th column instead of 5th one).

Result of search on 800Kbp subtelomeres with parameters as above may be found at `data/hg002/repeats_800k`. High scores towards proximal ends of subtelomeres might suggest that setting subtelomere length to 500Kbp might not be enough.

## References

- [1] Rhie A, et al. The complete sequence of a human Y chromosome. *Nature*, 2023. Can be downloaded at `s3-us-west-2.amazonaws.com/human-pangenomics/T2T/HG002/assemblies/hg002v1.1.fasta.gz`
- [2] Wei Shen, Botond Sipos, and Liuyang Zhao. 2024. SeqKit2: A Swiss Army Knife for Sequence and Alignment Processing. *iMeta* e191. doi:10.1002/imt2.191.
- [3] Quinlan AR and Hall IM, 2010. BEDTools: a flexible suite of utilities for comparing genomic features. *Bioinformatics*. 26, 6, pp. 841–842.
- [4] Seqtk. <https://github.com/lh3/seqtk>
- [5] ImageMagick. <https://imagemagick.org/index.php>
- [6] Alexander P Sweeten, Michael C Schatz, Adam M Phillippy, ModDotPlot—rapid and interactive visualization of tandem repeats, *Bioinformatics*, Volume 40, Issue 8, August 2024, btae493, <https://doi.org/10.1093/bioinformatics/btae493>
- [7] Sievers F, Barton GJ, Higgins DG (2020) Multiple Sequence Alignment. *Bioinformatics* 227, pp 227-250, AD Baxevanis, GD Bader, DS Wishart (Eds)